# A Republic, Not a Democracy

## The Legacy of John Birch Society Membership

Hong Wang

Boston University
Department of Economics
hongxi@bu.edu

July 6, 2026

## Abstract

Economic decline and demographic threat dominate explanations of contemporary right-wing mobilization, but older right-wing organizations are rarely measured directly. The John Birch Society (JBS) offers a rare window onto durable pre-MAGA right-wing infrastructure: newly digitized rolls from a secretive Cold War conservative organization. I recover 45,188 member records from archival dot-matrix printouts and convert them into a county-level measure of 1980s JBS presence. Within the replicated January 6 specification, counties with more JBS members produced more DOJ-charged defendants, and including JBS presence attenuates the manufacturing-decline association while leaving white-population decline intact. Yet JBS-rich counties show less Trump movement and no link to diffuse right-wing terrorism. The cross-outcome pattern points to older organizational infrastructure rather than generic grievance, partisan conversion, or right-wing violence of every kind.

**JEL Codes:** N42; D72; D74; C81.

**Keywords:** John Birch Society; organizational persistence; political violence; digitization; OCR.

**Acknowledgments:** I thank Melissa Dell, Martin Fiszbein, and Robert Margo for their encouragement and advice. I thank Ryan (Muhammad) Adriansyah and Daniel Dugan for their excellent research assistance in digitizing historical documents. Special thanks to Jasmine Sykes-Kunk and other archivists at the Brown University Library for helping me access the John Birch Society Collection, the BU Institute for Economic Development for funding part of this digitization, and the BU Shared Computing Cluster for providing the GPU resources.

Membership rolls and administrative records can reveal the local networks behind political outcomes, but many survive only as physical pages, scanned images, or archival records. For contemporary right-wing mobilization, this archival gap matters because recent economic and demographic shocks are easier to observe than the older organizations that connected those conditions to political action. Without direct measures of those networks, researchers cannot easily distinguish new grievance from long-standing mobilizing capacity.

I introduce a new geocoded dataset of John Birch Society (JBS) members in the 1980s. Founded in 1958 by a group of wealthy businessmen, the JBS grew into a nationwide network of secretive local chapters that shaped decades of conservative mobilization (Dallek 2023). JBS members and their organizational tactics helped Phyllis Schlafly defeat the Equal Rights Amendment (Dallek 2023; Spruill 2017), and the organization's ideological legacy—anti-communism, free-market fundamentalism, conspiratorial thinking, and the insistence that America is "a republic, not a democracy"—continues to reverberate through contemporary conservative politics. Despite this influence, no prior study has digitized the complete JBS membership rolls at the individual-record level or geocoded the full population of members.

This paper contributes to a literature on the long-run persistence of political behavior and the organizational channels that carry it. Local political cultures can persist across generations (Voigtländer and Voth 2012; Acharya, Blackwell, and Sen 2016). Closest to my design, Satyanath, Voigtländer, and Voth (2017) use hand-collected rosters of interwar German civic associations to show that denser associational networks accelerated entry into the Nazi Party: organizational infrastructure, not grievance alone, converted discontent into mobilization. For the United States, Fryer and Levitt (2012) digitized 1920s Klan rosters to study membership and its consequences, McVeigh, Cunningham, and Farrell (2014) show that counties with 1960s Klan activism realigned toward the Republican Party over subsequent decades, and Madestam et al. (2013) find that even brief Tea Party mobilization left durable organizational residue. Relative to this work, the JBS rolls offer something rare: a comprehensive individual-record census of a national postwar right-wing membership organization, measured at fine geography, for an organization whose networks survived into the contemporary period. The paper extends the persistence literature from historical shocks and defunct organizations to infrastructure directly upstream of present-day politics.

The data come from the JBS's comprehensive 1987 membership rolls: 45,188 member records recovered from dot-matrix printout pages. Unlike the mailing-list data used by Kraft (1992), these alphabetical rolls preserve internal chapter markers, allowing me to distinguish local chapter members from members of the national "HOME" chapter while aggregating records to counties and congressional districts.

I demonstrate the dataset's value by adding county-level JBS membership to three existing studies of contemporary right-wing politics: Pape, Larson, and Ruby (2024) on January 6 defendants, Broz, Frieden, and Weymouth (2021) on Trump's populist vote gains, and Nemeth and Hansen (2022) on right-wing terrorism. The applications do not estimate the causal effect of joining or hosting the JBS. Instead, I use a pre-existing organizational stock to test whether modern outcomes align with older conservative infrastructure, and whether published county-level relationships change when I measure that infrastructure directly. The results are strongest for January 6. Counties with more JBS members in the 1980s were significantly more likely to be home counties of DOJ-charged January 6 defendants. Adding JBS membership also changes the economic-grievance result in prior models: manufacturing decline becomes null, while the central white-population-decline result remains intact. The same JBS-rich counties, however, show less 2012–2016 electoral movement toward Trump and no systematic relationship with diffuse right-wing terrorism. The pattern is more consistent with older organizational infrastructure than with generic grievance, Trump-era electoral conversion, or right-wing violence of every kind. Figure 1 summarizes the cross-outcome pattern.

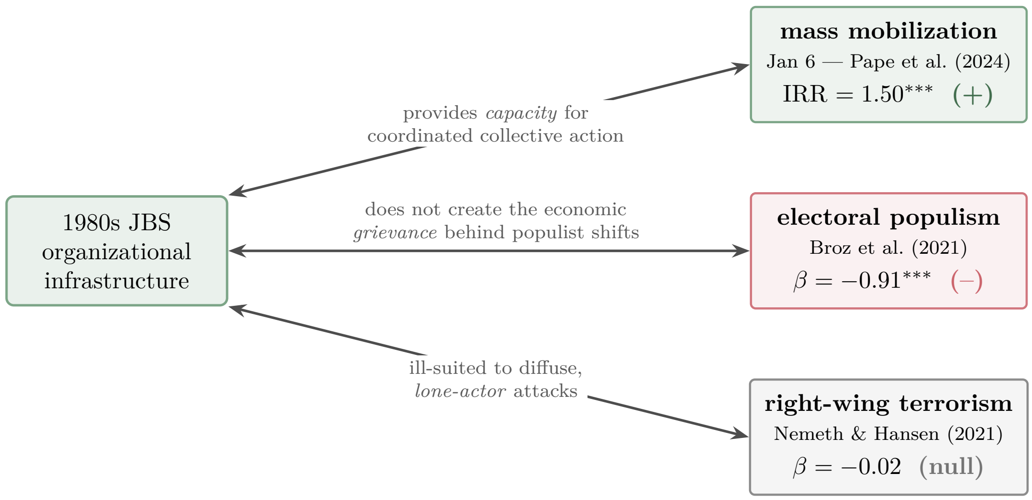


Figure 1: Conceptual summary of the three augmented replications. The same organizational stock is associated with coordinated mobilization, negatively associated with electoral conversion, and unrelated to diffuse terrorism, consistent with infrastructure operating as capacity, not grievance. Double-headed arrows denote descriptive association, not causation.

The research bottleneck for projects like this one is the cost and difficulty of converting unstructured page images into structured, analyzable data. Fryer and Levitt (2012) digitized Ku Klux Klan membership rolls from multiple archives to study who joined the 1920s Klan and its social and political impact.[1] Dell et al. (2023) demonstrate the potential of custom deep learning at massive scale, digitizing over 430 million historical newspaper articles with a modular pipeline. Dell (2025) makes a comprehensive case that custom pipelines consistently outperform both traditional methods and off-the-shelf commercial tools for extracting structured data from unstructured text and images, especially when historical sources differ sharply from modern web text.

General-purpose LLM APIs offer an appealing shortcut: send page images to a hosted model, receive fluent transcriptions, and report a character error rate estimated on a random validation sample. For research data, that error model is too thin in three ways. First, an aggregate character error rate treats errors as exchangeable, but a one-character error in a ZIP code relocates a member to a different county, while the same error in a surname is harmless for spatial analysis. Second, LLM transcription errors tend to be confabulations—well-formed records with plausible but wrong fields—so they pass both casual inspection and validity screens; a fabricated ZIP code can still be a valid ZIP code. Third, a sampled

[1] More recent work has compiled larger collections of KKK rosters to examine the relationship between military service and Klan participation.

error rate cannot reveal whether mistakes are systematic—concentrated on faint pages, particular record formats, or particular regions—which is precisely the structure that converts transcription noise into spatially correlated measurement error in downstream regressions (Dell 2025). The pipeline below is designed around the opposite standard: errors should be systematic, flaggable, and traceable to source images, so verification bounds them rather than merely estimating them.

Recovering the JBS rolls required a custom OCR digitization and verification pipeline. I use a modular deep learning pipeline for layout detection, line segmentation, optical character recognition, parsing, and geocoding. The pipeline achieves a final character error rate of 0.32%. More importantly, its error structure is auditable: research assistants verified lines against source images, and validation scripts screened geographic fields using ZIP-code, state, city, and alphabetical-order constraints.

The paper makes three contributions. First, I introduce the first individual-record, geocoded dataset of comprehensive JBS membership rolls, making it possible to measure conservative organizational infrastructure directly. Second, I show how that measure changes the interpretation of recent county-level findings on populism and political violence. Third, I provide a practical guide to digitizing structured archival membership records, including the pipeline's architectural choices, failure modes, and workarounds when off-the-shelf OCR or LLM API calls are difficult to audit.

The paper proceeds as follows. Section 1 provides historical background on the John Birch Society and its connection to the anti-ERA movement. Section 2 describes the archival membership rolls. Section 3 presents the digitization and verification pipeline. Section 4 describes the geography of JBS membership. Section 5 uses the dataset to revisit existing studies of contemporary right-wing politics. Section 6 concludes by discussing interpretation, limitations, and future applications of the data.

# 1 The Birchers as Organizational Infrastructure

The John Birch Society built a national infrastructure of local chapters, recurring meetings, internal directives, front groups, newsletters, and overlapping conservative social networks. Its membership rolls offer a county-level map of conservative organizational infrastructure: the places where a secretive but nationally coordinated right-wing organization had members, routines, and potential mobilizing capacity.

Robert Welch Jr., a retired candy manufacturer from a family of Baptist preachers, founded the JBS in 1958 with several wealthy businessmen, including Fred Koch of Koch Industries, Robert Waring Stoddard of Wyman-Gordon, and Harry Lynde Bradley of the Allen-Bradley Company. The organization took its name from John Birch, a Baptist missionary and U.S. military intelligence officer killed by Chinese Communist forces in 1945, whom Welch regarded as the first American casualty of the Cold War. Welch envisioned the JBS as "humanity's last hope" against what he believed was an imminent Communist takeover of the United States, orchestrated by a shadowy cabal of elites pursuing a New World Order.

The JBS combined free-market ideology with conspiratorial anti-communism in ways that distinguished it from mainstream conservatism. The JBS promoted the view that Republican President Dwight Eisenhower was part of the Communist conspiracy, that fluoridation of public water was a Communist plot to control Americans, and that the Civil Rights Movement and Martin Luther King Jr. were instruments of Soviet subversion.[2] The JBS also popularized the claim that America is “a republic, not a democracy,” an idea that continues to appear in contemporary conservative discourse.

The organization’s tactics were as important as its ideology. Welch, drawing on his background in marketing, understood the importance of packaging political ideas for mass consumption. The JBS maintained a prescriptive monthly bulletin and published *American Opinion*, through which national leadership directed the activities of local chapters. The organization expected members to attend monthly meetings, though enforcement varied. The JBS capped local chapters at twenty members to preserve secrecy, a structure explicitly modeled on small Communist cells. The JBS also created front organizations with innocuous names, allowing members to pursue organizational goals without the stigma of direct association. Chapter routines, front groups, and national directives make JBS membership a plausible measure of local organizational capacity.

The JBS initially drew members from the upper echelons of American society, including wealthy businessmen and their social circles, before expanding to white upper-middle-class professionals such as doctors and dentists.[3] Members paid dues that varied by gender: the JBS charged men $24 per year, approximately $260 in 2024 dollars, while women paid half that amount. Mainstream conservative figures attempted to distance the movement from the JBS—William F. Buckley Jr.’s *National Review* publicly denounced its more extreme claims—yet these denunciations appear to have deepened the commitment of existing members (Dallek 2023).[4]

The JBS was especially active in suburban conservative mobilization. In Pasadena and Orange County, the organization opposed school integration while also rallying economic elites against labor unions and the welfare state (McGirr 2001). Together, cultural conservatism, economic libertarianism, and grassroots suburban organizing anticipated strategies that would later become central to the modern Republican coalition. The key implication for this paper is that JBS membership reflected local social and political networks as well as national ideological discourse.

Phyllis Schlafly’s campaign against the Equal Rights Amendment illustrates the connection between JBS networks and broader conservative mobilization. Archival correspondence and JBS internal records indicate that Schlafly belonged to the John Birch Society, though she later denied it, and that she resigned strategically in 1964 to prevent *A Choice Not an Echo* from being associated with the organization during the Goldwater campaign.[5] Dallek

[2] The fluoridation conspiracy was satirized in Stanley Kubrick’s 1964 black comedy *Dr. Strangelove*, where General Jack D. Ripper warns that fluoridation threatens Americans’ “precious bodily fluids.”

[3] Kraft (1992) found that physicians and dentists were disproportionately represented in the JBS mailing list: 3.54% of his sample held the title “Dr” or “MD,” compared to approximately 0.25% of the general population who were physicians.

[4] Archival documents show that the JBS revoked the membership of individuals it found to be associated with the KKK.

[5] In a December 5, 1959 letter to Verne Kaub of the American Council of Christian Laymen, Schlafly wrote that she and her husband “both joined promptly after the Chicago meeting.” The February 1960 *JBS Bulletin* described her as “a very loyal member.” In her oral history, however, Schlafly said only that she had attended “a seminar that Robert Welch spoke at in Chicago” (DePue 2011). In an August 19, 1964 letter

(2023) documents JBS members leaving the official rosters while maintaining close informal ties to the organization.

Through STOP ERA, Eagle Forum, and *The Phyllis Schlafly Report*, Schlafly united evangelical Protestant and Catholic women against ERA ratification. She personally selected state chairwomen for STOP ERA, drawing on overlapping contacts from the National Federation of Republican Women, the National Council of Catholic Women, the Daughters of the American Revolution, the Goldwater campaign, and other conservative organizations (Critchlow 2005). Several state chairs drew on JBS networks to build their operations, although the extent to which Schlafly herself directed this is not documented. Dorothy (Dot) Malone Slade, a JBS member who knew Schlafly through the NFRW, led North Carolina and invited fellow JBS members to help. Beverly Findley, a former JBS member, led Oklahoma with Ann Patterson and worked with JBS members in planning anti-ERA hearings. In Mississippi and Oklahoma, anti-ERA organizing drew on the State Farm Bureau, which had close institutional ties to the JBS. In Utah, JBS women organized against the ERA through a front group called HOTDOG—Humanitarians Opposed to Degrading Our Girls (Spruill 2017).[6]

Because the JBS operated through local chapters, front groups, newsletters, conservative women's organizations, and informal social ties, the membership rolls provide a rare county-level measure of where that infrastructure had taken root.

# 2 The 1987 Membership Rolls

The project data comes from the John Birch Society records held at the John Hay Library at Brown University. When the JBS relocated its headquarters from Belmont, Massachusetts, to Appleton, Wisconsin, in 1989, Chip Berlet of Political Research Associates (PRA) recovered discarded headquarters materials from a dumpster outside the Belmont office; PRA later donated those materials to Brown (Kraft 1992). The Brown materials include several membership-related sources. Kraft (1992) analyzed a JBS mailing list, then held at PRA and labeled JMS 344, which included first-class postage, second-class postage, and complimentary-list sections.[7]

The source digitized here is distinct from the Kraft 1992 mailing list.[8] I digitize the JBS's alphabetical membership rolls, which organize members by surname, not postal class. The rolls appear to be a comprehensive internal membership record as of 1987 and contain 45,188 records across 1,870 scanned page images of continuous-feed dot-matrix printout.[9] In total, the rolls contain 99,183 text lines. Appendix Table 3 reports 638 joint spousal records, usually

to Mrs. Tom Anderson, a JBS National Council member's wife, Welch confirmed that the decision not to promote the book through JBS bookstores was at Schlafly's request (Lazar 2020).

[6]In 1977, JBS leader John F. McManus claimed the JBS had members "in leadership positions in every state where the ERA was scotched" (Spruill 2017).

[7]The second-class postage entries begin with foreign addresses before proceeding to domestic members, corresponding to the categories described in his study. Excluding the complimentary list, Kraft counted 21,305 individual and joint spousal memberships and 58 foreign members.

[8]I did not digitize the mailing-list sections analyzed by Kraft 1992.

[9]The page count refers to scanned PDF page images, not to the printout's internal page numbers, which do not correspond consistently to the alphabetical sequence of entries.

entries beginning with "Mr & Mrs." The dataset therefore counts membership records, not unique individuals. Unlike state totals or sampled ZIP-code data, the rolls preserve individual records and internal chapter markers while supporting county- and district-level aggregation. In Section 4, I use that structure to measure conservative organizational presence before the contemporary populist period.

# 3 Digitization and Verification

The central digitization challenge was spatial accuracy. The empirical analyses depend on county and district assignment, so an OCR error in a ZIP code, city, or state can move a member to the wrong place. I therefore designed the pipeline around auditable accuracy: each stage produces intermediate outputs that reviewers can inspect, correct, and link back to the source image.

The pipeline has four stages. First, layout detection identifies the region of each page that contains membership records. Second, line detection segments that region into individual text-line images. Third, OCR reads the content of each line. Fourth, post-processing parses the OCR output into structured member records and assigns geographic units. Figure 2 summarizes the workflow.

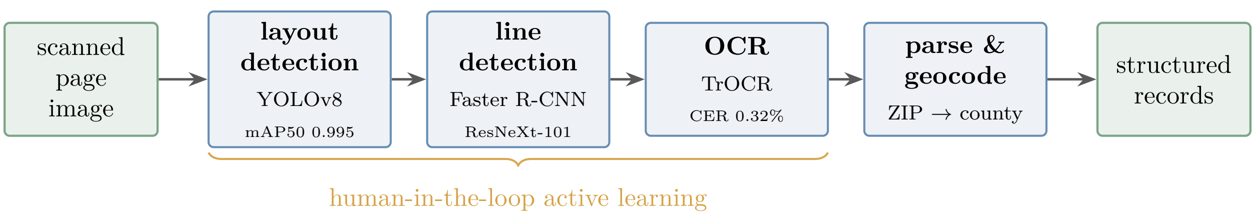


Figure 2: Pipeline architecture and division of labor. The brace marks the stages where active learning lowers marginal labeling costs; parsing and geocoding are deterministic post-processing steps, not learned predictions.

## 3.1 Layout Detection

Layout detection separates the membership-record area from headers, page numbers, margins, and other non-data material. Although the relevant region is visually obvious to a human, isolating it prevents later stages from treating page artifacts as member records. I treated layout detection as a human-in-the-loop task, using Label Studio (Tkachenko et al. 2020) with a fine-tuned small YOLOv8 model (Jocher, Chaurasia, and Qiu 2023). Initial hand labels generated model predictions; I folded corrected predictions back into training.

The layout stage required labeling 1,460 of 1,870 page images. The final model achieved mAP50 of 0.995 and mAP50–95 of 0.962, indicating near-perfect recovery of the data region. Because the target region is large and visually regular, a small YOLOv8 model trained with default hyperparameters was sufficient.

## 3.2 Line Detection

Line detection was the most technically demanding stage. Each page contains roughly 53 text lines, and each line is long and narrow, approximately $16 \times 659$ pixels. The line geometry is poorly suited to object detectors optimized for compact natural-image objects. A YOLOv8 model, successful in the layout stage, performed poorly for line detection, reaching only 47% mean average precision.

The failure reflected box geometry, not computation alone. Text lines require proposal boxes that are many times wider than they are tall. I therefore switched to Detectron2 (Wu et al. 2019) with a Faster R-CNN architecture and a ResNeXt-101 backbone pretrained on COCO (Lin et al. 2014). I customized the anchor ratios to 0.02–0.04, corresponding to boxes 25 to 50 times wider than tall. Figure 3 illustrates why this change was necessary.

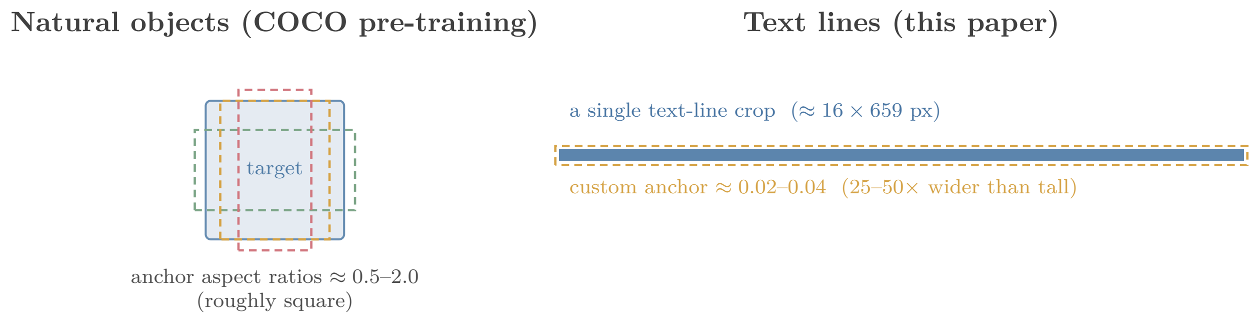


Figure 3: Anchor geometry, not model size, is the binding design choice for line detection. Natural-image pretraining supplies useful visual features, but the proposal network must be re-tuned to the long, thin boxes that membership-roll text lines occupy.

I trained the line-detection model on 7,732 page crops derived from the layout labels. Over training, total loss fell from 2.18 to 0.17, RPN classification loss fell from 0.686 to below 0.001, and foreground classification accuracy rose from 14.3% to 96.7%. I exported model predictions directly into the same Label Studio project where reviewers labeled and corrected line boxes. The export workflow kept prediction and correction in the same labeling environment.

## 3.3 Optical Character Recognition

After line segmentation, OCR reads the text from each line image. I fine-tuned the TrOCR model (Li et al. 2023), a transformer-based OCR architecture, on the visual characteristics of the JBS dot-matrix printout. Early tests with generic OCR systems, including Apple OCR, helped bootstrap initial labels but produced too many errors for final data construction. Those predictions also revealed the source's characteristic error structure: visually similar dot-matrix characters, dropped punctuation, and inconsistent spacing.

At all stages, I used an active-learning workflow. Research assistants corrected batches of line images; I added the corrected lines to the training set; the updated model generated predictions for the next batch. As accuracy improved, each batch required fewer corrections. Human review remained constant, but the nature of the task shifted as the amount of human typing declined sharply over time.

The OCR training set advanced sequentially through the alphabet. Random sampling would better estimate population-level error, but the goal was full-corpus digitization; the fixed record format and visual features do not vary systematically by surname.

The final OCR training and validation set contains 82,797 corrected text lines, with 74,517 used for training and 8,280 for validation. After 58 passes over the training data, the model achieved a validation character error rate of 0.32%, or fewer than one incorrect character per 300 characters on average. The final model processed the remaining 16,386 line crops, and research assistants verified those lines against the source images.[10] Table 1 reports the decline in character error as the training set expanded.

Table 1: OCR active-learning returns as labeled lines increase

| **Labeled Lines** | **Corpus %** | **Validation CER %** |
|---:|:---:|:---:|
| 2,415 | 2.4% | 2.95 |
| 5,889 | 5.9% | 1.85 |
| 9,600 | 9.7% | 1.19 |
| 82,797 | 83.5% | 0.32 |

*Notes:* CER is measured on a 20% held-out validation set. The total corpus contains 99,183 lines.

## 3.4 Post-Processing and Geocoding

The raw OCR output is a sequence of text lines. I link those lines into member records and parse them into fields. Each JBS record generally spans two lines: the first contains the surname, given name, household code, salutation, and street address; the second contains the city, state, ZIP code, chapter marker, and administrative codes. I infer gender from salutations such as Mr., Mrs., Ms., and Miss. Because the rolls use a regular fixed-position format, parsing depends on reproducibility and field preservation. The parser preserves tokens in a structured schema, retains administrative codes, and assigns counties using ZIP Code Tabulation Area (ZCTA) crosswalks. Figure 4 shows a representative record parse.

## 3.5 Verification and Error Auditing

The verification design has two layers. First, research assistants checked every text line against the source image. Hence, no line entered the final dataset without human review; each parsed field therefore traces back to a line that a research assistant typed or checked against the original image.

Second, the pipeline exploits the structure of the records to catch residual errors. The most common OCR confusions involve visually similar characters in dot-matrix print, especially 0/O, 1/I, and 5/S. Dot-matrix confusions can corrupt names and administrative fields,

[10] Additional training details for all pipeline stages, including learning rate, batch size, anchor sizes, and evaluation schedule, appear in Appendix Table 2.

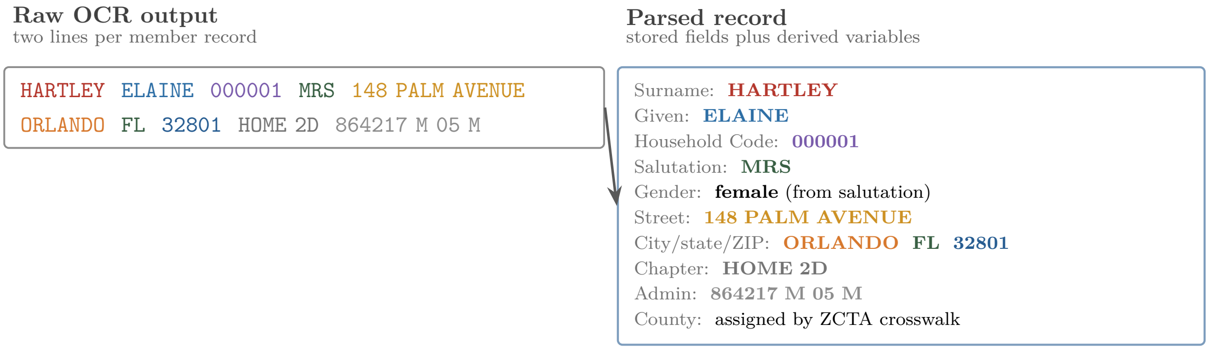


Figure 4: Fixed-position parsing converts two OCR lines into analyzable fields. The anonymized example uses synthetic personal identifiers while preserving the record structure. The parser derives gender from salutations, retains the household code field, keeps chapter markers distinct from administrative codes, and assigns county after ZIP parsing through a ZCTA crosswalk. The example record is fictitious and used only for illustration.

but geographic errors are most consequential for this paper. Validation scripts checked ZIP codes against valid entries, state abbreviations against a fixed list, and city–ZIP combinations against known geographic matches, then flagged failed records for manual review. The alphabetical ordering of the rolls provides a final structural check. Because pages proceed by surname, any record outside the expected alphabetical range of its page indicates a linking, parsing, or OCR problem.

The final dataset passes this check with no breaks across all 1,870 pages. The verification strategy makes county- and district-level analysis possible. The relevant claim is that the pipeline made errors visible, bounded, and correctable, producing a set of geocoded membership records with spatial assignments traceable to verified source lines.

# 4 The Geography of JBS Membership

The verified rolls contain 45,188 geocoded member records, of which the county-level analysis covers 2,456 U.S. counties in all 50 states.[11] Because some records list married couples jointly, the number of individuals represented is somewhat higher than the number of records. The county distribution is highly skewed: the median county had four members, while Los Angeles County alone had 1,582. Overall, the JBS had at least one recorded member in 78% of U.S. counties.

[11] Of 45,188 records, 46 geocode to areas outside the county analysis universe: San Juan, Puerto Rico; St. Croix, U.S. Virgin Islands; and the Chugach and Copper River census areas created by the 2019 split of Alaska's Valdez–Cordova area. County-level analyses therefore cover 2,456 counties.

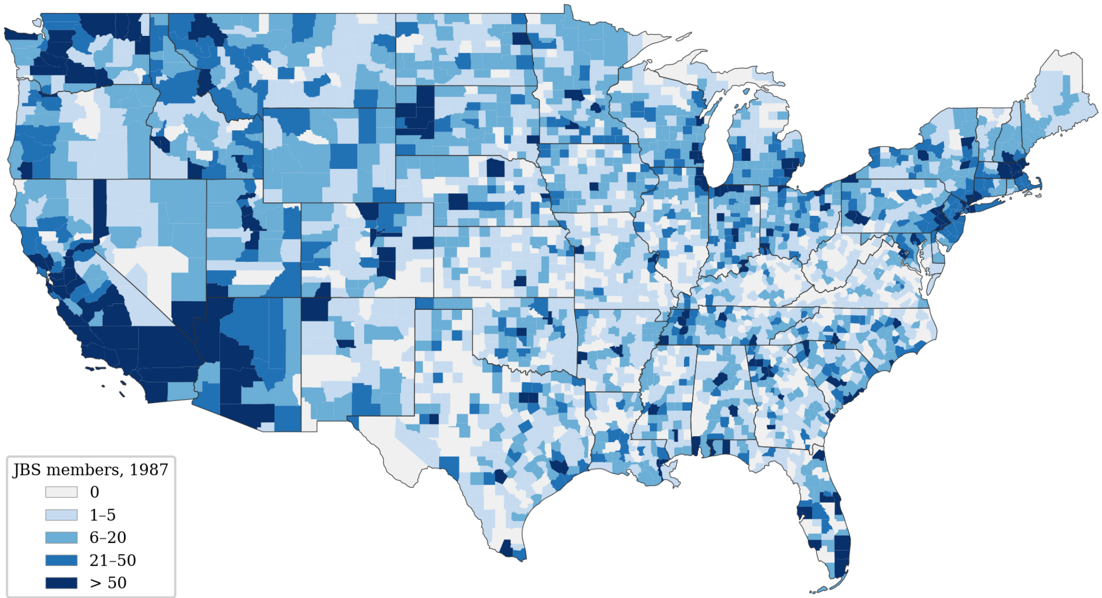


Figure 5: County-level geography of JBS membership in the 1987 rolls. The map displays the historical organizational stock used in the empirical applications below.

## 4.1 Geographic Distribution

The absolute geography of membership reflects the organization's known strongholds. California had the largest number of records, consistent with historical accounts of JBS organizing in Orange County and the Los Angeles suburbs (Dallek 2023; McGirr 2001). Texas and Florida ranked second and third. Per capita membership was highest in South Dakota, Idaho, Montana, and North Dakota.

The rolls map organizational presence, not conservative population alone. JBS members received communications, could attend local meetings, and could connect through chapters, front groups, and overlapping conservative networks. The measure therefore captures a historical stock of organizational infrastructure, not a contemporaneous vote outcome.

Appendix Table 3 summarizes membership by gender. Men outnumbered women roughly 3:2, consistent with the JBS's origins in male business networks. The substantial female membership also fits historical accounts of women sustaining local JBS organizing even as men dominated much of the national leadership (Dallek 2023).

## 4.2 Validation Against Prior Work

The geography of the digitized rolls closely matches the state-level patterns reported by Kraft (1992), even though the underlying source differs. Kraft analyzed a smaller active mailing list and found that California, Texas, and Florida had the largest absolute memberships, with 2,648, 1,232, and 1,154 members, respectively. In the alphabetical rolls digitized here, the same three states rank at the top: California has 6,274 records, Texas has 3,048, and Florida has 2,190. Across states, the Spearman rank correlation between Kraft's counts and the digitized rolls is $\rho = 0.93$ ($p < 10^{-20}$). The high rank correlation provides external validation for the recovered geography. The rolls are larger than Kraft's mailing list, but they preserve the same broad spatial distribution.

### 4.3 Distinguishing the 1987 JBS from 1964 Goldwater Geography

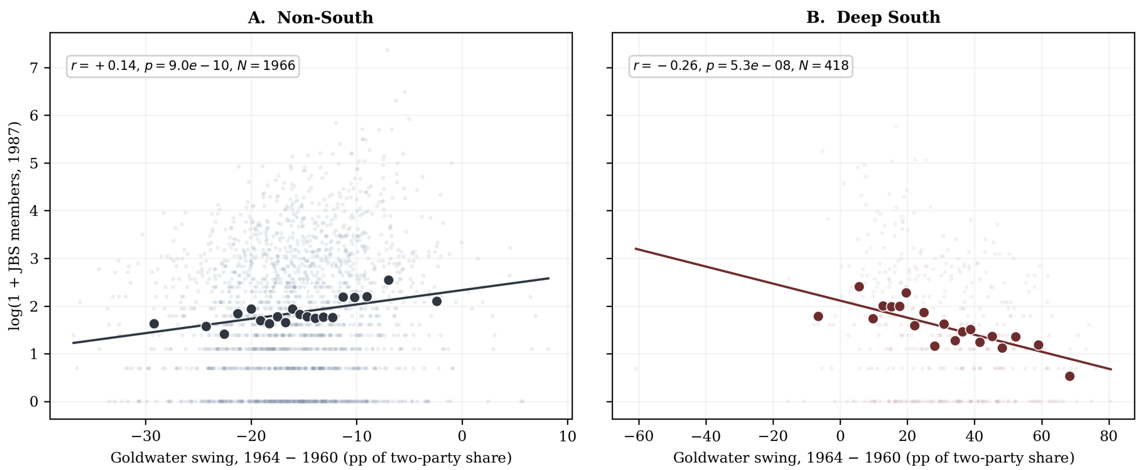


Figure 6: Regional decomposition of the county-level relationship between the Goldwater swing (Republican two-party share in 1964 minus 1960) and 1987 JBS membership. Binned means (20 quantile bins) overlay the county scatter. Outside the South the swing is positively associated with later JBS membership ($r = 0.14$, $p < 10^{-9}$); in the Deep South the relationship is negative ($r = -0.26$, $p < 10^{-7}$).

A natural concern is that JBS membership may reproduce the geography of Goldwater-era conservative mobilization, especially in Southern California (Dallek 2023; McGirr 2001). County-level presidential returns suggest otherwise. I measure Goldwater-era mobilization as the swing in the Republican two-party share from Nixon in 1960 to Goldwater in 1964. Nationally, the relationship between the Goldwater swing and 1987 JBS membership is weak ($r = -0.07$) because two much stronger regional patterns pull in opposite directions. Outside the South, the swing is positively correlated with later JBS membership ($r = 0.14$, $p < 10^{-9}$), consistent with a free-market, anti-communist conservative tradition. In the Deep South, the relationship is negative ($r = -0.26$, $p < 10^{-7}$),[12] reflecting a different regional coalition organized around civil-rights backlash and rural Black Belt politics.

The regional decomposition supports the interpretation of the JBS rolls as a specific measure of conservative organizational infrastructure, not generic conservative partisanship. Even in Georgia, where both Goldwater support and JBS organizing were substantial, JBS membership clustered in suburban professional counties, while the Deep South Goldwater pattern was more rural and Black Belt. The JBS measure therefore captures one strand of the conservative coalition: organized, anti-communist, free-market, suburban conservatism. Figure 6 summarizes this regional decomposition.

[12] The Deep South pattern is not driven by Alabama's unusual 1964 ballot context: excluding Alabama yields a similar negative correlation between the Goldwater swing and later JBS membership.

# 5 Applications: Organizational Infrastructure and Contemporary Right-Wing Politics

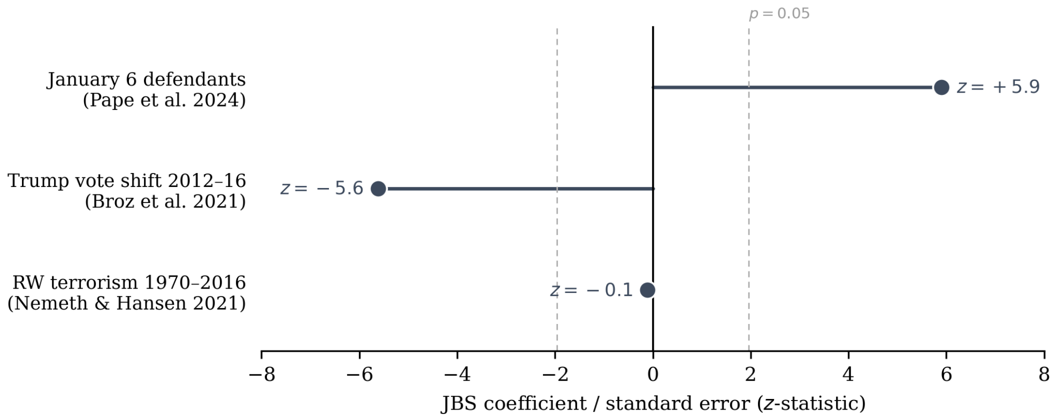


Figure 7: Direction and precision of the standardized JBS coefficient in each replication model.

*Notes:* Each bar reports the JBS coefficient divided by its standard error; dashed lines mark $|z| = 1.96$. Native-unit estimates are log-IRR 0.408 (SE 0.069) for January 6 defendants, $-0.912$ percentage points (SE 0.162) for the 2012–2016 Trump vote shift, and logit $-0.020$ (SE 0.141) for right-wing terrorism. JBS membership is measured as $\log(1 + \text{county member count})$, standardized to mean zero and unit variance within each estimation sample, with all original controls.

I use the JBS rolls to revisit three recent county-level studies of right-wing politics: Pape, Larson, and Ruby (2024) on January 6 defendants, Broz, Frieden, and Weymouth (2021) on Trump's 2016 vote gains, and Nemeth and Hansen (2022) on right-wing terrorism. I selected these studies because they examine related but distinct outcomes—January 6 participation, electoral conversion, and diffuse right-wing terrorism—and because their replication files make it possible to add a common historical measure across models. I also revisit Pape, Larson, and Ruby (2024)'s 2020 election-certification split because it is the one elite-level condition in their analysis that amplified the January 6 pattern. Figure 7 summarizes these results across models.

In each application, I measure JBS membership as $\log(1 + \text{county member count})$, standardized to mean zero and unit variance within each estimation sample. I assign members to counties using the same ZCTA crosswalks. All negative binomial models report HC1 robust standard errors, following Pape, Larson, and Ruby (2024); the terrorism logits cluster standard errors by county. JBS membership was not randomly assigned to counties, and places with JBS members likely differed from other places in ideology, social networks, local elite structure, and anti-government sentiment. The exercise therefore does not estimate the causal effect of JBS membership itself. Instead, it tests whether a descriptive, direct measure of pre-existing conservative organizational infrastructure changes the interpretation of published county-level relationships.

## 5.1 Mobilizing Populism: January 6 Capitol Defendants

Pape, Larson, and Ruby (2024) analyze county-level predictors of individuals charged by the Department of Justice for the January 6 Capitol breach. Their main finding is that white population decline from 2010 to 2020 is the strongest county-level predictor of January 6 defendants. Manufacturing decline from 1970 to 2020 is also positive, but smaller, leading the authors to emphasize demographic threat over economic grievance.

The descriptive overlap is already large: 415 of 458 defendant counties (90.6%) had at least one JBS member in the 1987 rolls, compared with 78% of all U.S. counties. Among counties with any JBS members, counties above the JBS median averaged 0.645 defendants, compared with 0.082 in counties with no JBS members.

I replicate their full county model and then add JBS membership. The replication closely matches the published results: white population decline is the strongest predictor (incidence rate ratio [IRR][13] = 1.37, $p < 0.001$), manufacturing decline is positive and significant (IRR = 1.12, $p = 0.040$), and populist Trump support is negative (IRR = 0.77, $p < 0.001$). Adding JBS membership produces a large and precisely estimated association. A one-standard-deviation increase in logged JBS membership predicts 50% more January 6 defendants (IRR = 1.50, $p \approx 4 \times 10^{-9}$), net of population, urban status, distance to Washington, DC, and the original county controls. The model's AIC falls by 45.6 points, indicating substantially better fit even after penalizing the added JBS term.[14]

The JBS measure also changes the interpretation of the manufacturing result. Once JBS membership enters the model, manufacturing decline becomes null ($p = 0.040 \rightarrow p = 0.308$), while white population decline remains significant at a similar magnitude ($p = 1.3 \times 10^{-7}$). The pattern suggests that manufacturing decline partly marked counties where deindustrialization overlapped with older conservative organizational infrastructure. The result does not overturn the central result in Pape, Larson, and Ruby (2024): white population decline remains a distinct predictor. Rather, the added measure separates two components that are otherwise difficult to distinguish in county-level data: contemporary demographic threat and pre-existing mobilizing capacity. The association does not interact with demographic threat: a JBS $\times$ white-population-decline term is null (IRR = 1.03, $p = 0.55$), consistent with capacity and grievance operating additively rather than multiplicatively. The estimate is also insensitive to the most influential geography: dropping all of California (IRR = 1.47) or Los Angeles County alone (IRR = 1.47) leaves the coefficient essentially unchanged. Figure 8 visualizes the coefficient movement.

This result is nevertheless conditional on the published population functional form. Pape, Larson, and Ruby (2024) enter county population linearly, in units of 100,000 residents. Replacing that control with log population improves fit substantially (AIC 3,128.7 $\rightarrow$ 2,900.4) and attenuates every substantive covariate in the model, not only the JBS term: white population decline falls to IRR = 1.10 ($p = 0.048$), manufacturing decline and populist Trump

[13] Incidence rate ratios report multiplicative changes in expected counts. An IRR above 1 indicates a higher expected count; an IRR below 1 indicates a lower expected count. For example, an IRR of 1.50 means a 50 percent higher expected count for a one-standard-deviation increase in the standardized predictor, holding other covariates fixed.

[14] AIC compares model fit while penalizing additional parameters; lower values indicate better fit. Following Pape, Larson, and Ruby (2024), estimation excludes Alaska, which does not report presidential returns by county equivalent; $N$ = 3,112.

support become null, logged JBS membership falls to IRR = 0.93 ($p$ = 0.28), and a binary any-JBS indicator is also null under log-population control (IRR = 0.80, $p$ = 0.19). Defendant counts scale strongly with logged population (IRR = 4.67 in the binary-JBS specification), so the January 6 estimates should be read as conditional on the host specification. The cross-outcome contrast and the district-level certification differences below do not depend on this functional-form choice.

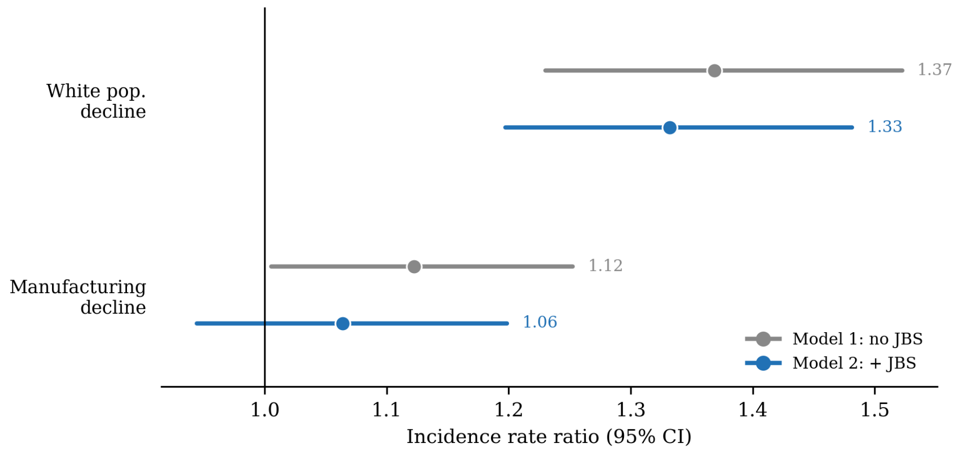


Figure 8: Manufacturing decline attenuates after adding observed JBS infrastructure, while white-population decline remains positive and precisely estimated. Points and 95% confidence intervals report incidence-rate ratios from the January 6 negative-binomial county models. The figure is descriptive and is not a causal mediation design.

## 5.2 Electoral Populism: Trump's Populist Vote Gains

The same JBS measure has a different relationship with electoral populism. Broz, Frieden, and Weymouth (2021) study Trump's county-level vote gains over Romney's 2012 performance. Their dependent variable captures the additional Republican two-party vote share Trump gained in 2016, not baseline Republican strength.

I replicate their county-level specification and add JBS membership. JBS membership is statistically significant, but the sign is negative ($\beta = -0.91$, SE = 0.16, $p < 10^{-7}$): counties with more JBS members showed less 2012–2016 shift toward Trump. Manufacturing decline remains significant in the baseline augmented model.

The negative coefficient does not contradict the January 6 finding, but the mechanism is not simple saturation. County-level mean reversion is real: Romney's 2012 two-party share strongly negatively predicts the 2012–2016 shift ($\beta = -0.15$, $p < 0.001$). Yet JBS membership is negatively correlated with the 2012 Republican baseline ($r = -0.24$)—the rolls concentrate in populous, metropolitan counties, not in the rural, high-baseline Republican counties—and the JBS coefficient survives the baseline control nearly intact ($-0.91 \rightarrow -0.76$, $p < 10^{-7}$). The direction also depends on the margin: an indicator for any JBS presence is weakly positive ($\beta = 0.49$, $p = 0.09$), while the intensity measure is strongly negative. The negative coefficient is robust to dropping California ($\beta = -0.78^{***}$) or Los Angeles County

($\beta = -0.94^{***}$), and to a per-capita specification ($\beta = -0.79^{***}$). The rolls therefore do not map the geography of Trump's incremental electoral breakthrough; they map an older, more metropolitan conservative infrastructure located in precisely the kinds of counties that shifted least toward Trump.[15]

## 5.3 Right-Wing Terrorism

The JBS measure does not predict all forms of right-wing violence. Nemeth and Hansen (2022) analyze county-year incidence of right-wing terrorist attacks from 1970 to 2016 and find that electoral competition in predominantly white counties is the strongest predictor. I replicate their main specification using pooled logit with county-clustered standard errors as an approximation of their mixed-effects logit and add JBS membership.

JBS membership is not statistically significant ($\beta = -0.02$, $p = 0.91$). The null result also holds in interactions with electoral competition ($p = 0.42$) and racial composition ($p = 0.26$). Restricting the panel to years after the membership snapshot does not change the conclusion: for 1988–2016, the coefficient is $\beta = -0.21$ ($p = 0.24$); for the pre-snapshot 1971–1987 window, it is $\beta = 0.24$ ($p = 0.34$). Readers should interpret this exercise cautiously: right-wing terrorism is rare in the data, with 192 events across 133,539 county-years, and the pooled logit does not reproduce the state-level random effects in the original model.

Even with those limitations, the null is informative. If JBS membership measured generic right-wing extremism, it should also predict diffuse right-wing terrorism. Instead, the contrast with January 6 suggests a narrower interpretation: the JBS was a political membership organization, not a paramilitary group, and its infrastructure was better suited to coordinated collective action than to low-frequency, often lone-actor violence.

## 5.4 Elite Contestation: 2020 Election Certification

As a supplementary application, I aggregate JBS membership to congressional districts and compare districts represented by House Republicans who voted to overturn the 2020 presidential election results with districts represented by members who certified the results. Districts represented by one of the 139 Republican objectors averaged 117.2 JBS members, compared with 93.3 in certifying districts (1.26×, one-sided Mann–Whitney $p = 0.002$). The association is similar by gender: male members are 1.26× higher in objector districts ($p = 0.001$), and female members are 1.22× higher ($p = 0.007$). By chapter type, the objector–certifier gap is larger for local chapter members (1.33×, $p = 0.005$) than for members of the national "HOME" chapter (1.14×, $p = 0.002$), but the chapter split should be read with care.

The two county-level series are highly correlated ($r = 0.65$), and in the January 6 county model the ordering does not replicate: entered jointly, HOME membership remains a strong predictor (IRR = 1.44) while the local-chapter coefficient shrinks toward the null (IRR = 1.11, $p = 0.15$). The 1987 vintage of the rolls complicates the interpretation of HOME membership, which by the late 1980s may reflect how members joined and remained affiliated

[15] Conditional on the 2012 baseline, manufacturing decline also attenuates ($p = 0.12$), consistent with Broz, Frieden, and Weymouth (2021)'s outcome capturing change relative to baseline.

as much as the absence of local organization. I therefore treat the chapter/HOME distinction as descriptive heterogeneity rather than as direct evidence on mechanism.

# 6 Conclusion

This paper introduced the first individual-record, geocoded dataset of John Birch Society membership and used it to test whether older conservative organizational infrastructure aligns with contemporary right-wing outcomes. The findings show why historical organizational measures matter for studies of contemporary right-wing politics. Recent shocks may help explain where grievances emerge, but older networks help explain where those grievances become organized political action.

The results are descriptive. JBS membership was not randomly assigned to counties, and the organization recruited in places where conservative ideology, local elite networks, and anti-government sentiment were already present. Readers should therefore interpret the JBS variable as a measure of pre-existing organizational infrastructure, not as the causal effect of JBS membership itself. Future work could strengthen causal identification by exploiting variation in the timing and geography of JBS organizing, including regional field staff, chapter coordinators, or organizing drives documented in the *JBS Bulletin*.

The timing of the measure also matters. The rolls capture a stock of members in 1987, not the full history of JBS activity from the organization's founding in 1958 onward. The 1987 timing makes the data especially useful for studying persistence into the contemporary period, but less well suited to explaining earlier outcomes in the 1960s and 1970s. Because the terrorism panel begins before 1987, the JBS rolls are better suited to detecting persistence after the membership snapshot than to explaining attacks earlier in the panel.

As Section 5 documents, the January 6 county coefficients—including the original study's headline predictors—are sensitive to the population functional form and should be read as conditional on the replicated specification.

The broader contribution is methodological as well as substantive. The digitization pipeline shows how researchers can convert structured archival records into auditable political data without treating OCR as a black box. The resulting dataset opens questions that were previously difficult to answer: how conservative organizational networks persisted across decades, how they intersected with anti-ERA mobilization and conservative women's organizations, how they shaped local Republican infrastructure, and how they influenced later conservative mobilization and election contestation. More generally, the findings suggest that studies of populism and political violence should pay closer attention to meso-level organization: the local networks and institutions that connect broad social conditions to political action.

# Data Availability and Ethics

The published paper redacts personal identifiers from source images. Public replication materials will release county and district aggregates, code, and crosswalks sufficient to reproduce the analyses, but not named individual-level membership records. Individual-level named data should be restricted to controlled replication use, subject to archival permissions and an IRB or institutional determination for secondary use of sensitive archival records.

# 7 Appendix

The appendix proceeds from data construction to empirical results. Sections 7.1–7.4 document archival access, scanning, model implementation, and verification. Sections 7.5–7.7 report supplementary descriptive statistics, full model estimates, robustness checks, and conceptual figures.

## 7.1 Archival Access and Scanning Workflow

The membership rolls consists of continuous-feed dot-matrix printout, making it physically legible but poorly suited to direct OCR. The pages are regular enough to support model-based extraction, but the combination of dot-matrix characters, faint print, page curvature, and long narrow text lines made off-the-shelf OCR unreliable.

Scanning required coordination with library staff because the records could not be removed from the archive, archive appointments had to be arranged well in advance, and scanner access was sometimes limited. The first collection pass took 20 days, from June 5 to June 25, 2024. The workflow unfolded over two years: research assistants corrected and verified lines during the 2024–25 academic year, and final validation and empirical integration continued through spring 2026. The conceptual framework—modular pipeline, fine-tuned detection, and human-in-the-loop labeling—came from Dell (2025)'s body of work. Dell's advice about matching object-detection anchors to target geometry shaped the line-detection stage: long text lines required radically different anchors than the natural objects these models are pretrained on.

## 7.2 Pipeline Implementation Details

Table 2 provides the model choices, training data, hyperparameters, metrics, and training time for each stage of the pipeline. Layout detection produced page-level crops, line detection produced line-level crops, OCR produced line-level predictions, and parsing produced structured records. Figure 9 shows representative source-image examples for the layout-detection, line-detection, and OCR stages of the workflow.

Table 2: Digitization pipeline implementation details

| Stage | Model | Training data | Hyperparameters, metric, and time |
|---|---|---|---|
| Layout Detection | YOLOv8n | 1,314 training crops (146 val.) | *Hyperparameters:* epochs: 100; optimizer: AdamW. *Metric:* mAP50: 99.5%; mAP50–95: 96%. *Time:* 10 days, including labeling and training. |
| Line Detection | Faster R-CNN (ResNeXt-101-32x8d-FPN) | 7,732 training crops | *Hyperparameters:* iterations: 224,700; lr: 1e-4; batch: 8; anchors: [.02-.04]. *Metric:* Classification accuracy: 96.7%. *Time:* 7 hours training. |
| OCR | TrOCR (base-printed) | 74,517 training lines (8,280 val.) | *Hyperparameters:* epochs: 60; lr: 4e-5; batch: 16; beams: 4. *Metric:* CER: 0.32%. *Time:* 21 hours training. |

*Notes:* All training was conducted on two NVIDIA GPUs with 40GB memory on Boston University's Shared Computing Cluster. Line segmentation used Detectron2 with a COCO-pretrained backbone. OCR used HuggingFace Transformers with Accelerate for multi-GPU training.

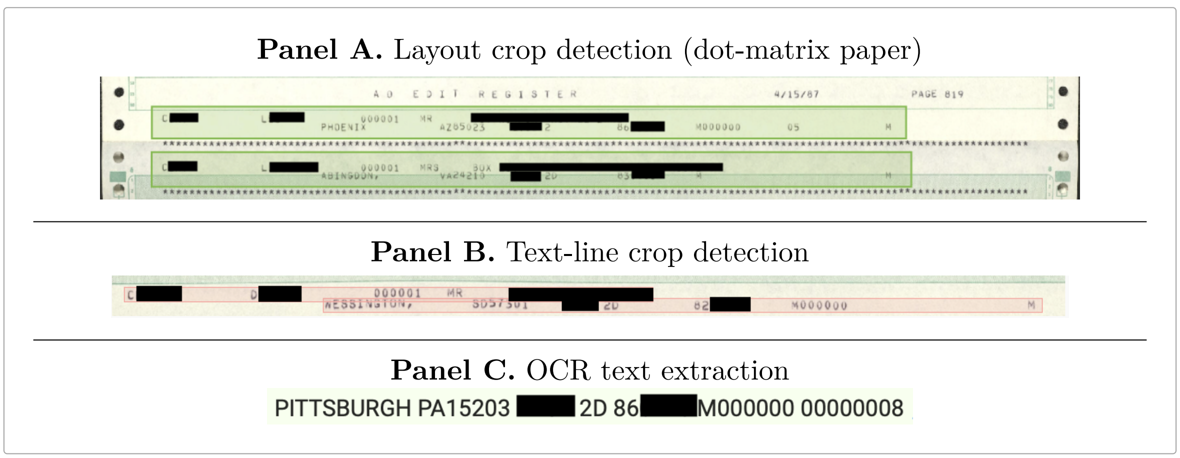


Figure 9: Representative redacted source-image examples for the digitization workflow. The panels show the visual inputs behind layout detection, line detection, and OCR, complementing the schematic pipeline in Figure 2. Black boxes indicate redactions of names, street addresses, and box numbers added for publication; city, state, ZIP code, chapter markers, and administrative codes are reproduced as they appear in the source scans.

## 7.3 Deskewing, Hashing, and Duplicate Handling

Several preprocessing experiments produced useful lessons. I initially tried deskewing because slight page rotation made line detection harder. Deskewing helped some images, but it also changed pixel coordinates and created bookkeeping problems across layout labels, line crops, and source-image links.

Image hashing—computing a perceptual fingerprint for each file—helped recover mappings when preprocessing changed file names or directory structure, but the added maintenance cost outweighed the accuracy gain. As the training set grew, the line detector became robust enough to residual skew that I stopped deskewing and rehashing as routine steps.

## 7.4 Research Assistant Workflow and Time Costs

The project required approximately 300 research-assistant hours across two research assistants over eight months. Manual transcription of 45,188 two-line records would have required substantially more labor, conservatively 1,500 to 2,000 hours at realistic transcription speeds. The largest bottleneck was line segmentation, because correcting bounding boxes around long narrow text strips requires pixel-level precision.

## 7.5 Supplementary Descriptive Tables

Table 3 reports membership records by salutation category and geographic coverage. Table 4 compares counties with JBS members, counties with January 6 defendants, and all counties. The tables supplement the descriptive discussion in Section 4.

Table 3: Membership records by salutation category and geographic coverage

| | Count | % |
|---|---|---|
| Male | 26,394 | 58.4 |
| Female | 16,385 | 36.3 |
| Couple | 638 | 1.4 |
| Professional | 1,271 | 2.8 |
| Other | 500 | 1.1 |
| Total records | 45,188 | 100.0 |
| Counties covered | 2,456 | of 3,141 |
| States covered | 50 | + DC |

*Notes:* Unit is a geocoded membership record, not necessarily a unique individual. Gender is inferred from salutation fields (Mr, Mrs, Ms, Miss, and related variants). "Professional" denotes records with MD, DDS, DVM, or similar title but no gendered salutation. "Couple" denotes joint spousal records (Mr & Mrs). Membership rolls date from mid-1987. County coverage excludes records geocoded outside the county analysis universe; see Section 4.

Table 4: Descriptive comparison of counties with JBS members and counties with January 6 defendants

| | JBS Counties ($n = 2{,}456$) | Defendant Counties ($n = 458$) | All Counties ($n = 3{,}141$) |
|---|---|---|---|
| *Key Independent Variables (means)* | | | |
| White pop. decline (%) | 4.42 | 5.62 | 4.15 |
| Mfg employment decline (%) | 8.25 | 12.60 | 8.05 |
| Trump populist support (%) | 4.89 | 1.29 | 5.40 |
| *Demographics* | | | |
| % Non-Hispanic white | 74.2 | 69.1 | 74.2 |
| Urban county (%) | 28.9 | 65.9 | 25.7 |
| *Overlap* | | | |
| Counties with both JBS & defendants | 415 of 458 defendant counties (90.6%) | | |
| Avg. defendants, high JBS | 0.645 vs. 0.082 in counties with no JBS | | |
| Spearman $\rho$ | 0.308 ($p < 10^{-70}$) | | |

*Notes:* JBS counties have $\geq 1$ member in the mid-1980s rolls. Defendant counties have $\geq 1$ individual charged by the DOJ for the January 6, 2021, Capitol breach as of April 2023. "High JBS" counties are above the median among counties with any JBS members.

## 7.6 Supplementary Empirical Tables

Table 5 reports the full January 6 negative binomial models, including the original replication, the model with JBS membership, a JBS-only control model, and the certification interaction. Table 8 summarizes the change in the manufacturing-decline coefficient before and after adding JBS membership. Tables 6–11 report replication fidelity, the outbidding specifications, and population-functional-form sensitivity. Table 9 reports the Broz et al. electoral-conversion specifications. Table 10 summarizes the terrorism logits and timing windows. Table 12 reports congressional-district differences by certification vote.

Table 5: Negative binomial models of county-level January 6 defendant counts, reported as incidence rate ratios

| | Model 1 Pape et al. Replication | Model 2 + JBS (log, z) | Model 3 JBS Only + Controls | Model 4 Certification Interaction |
|---|---|---|---|---|
| *Main Explanatory Variables* | | | | |
| White Population Decline 2010–2020 | 1.368*** | 1.332*** | | 1.334*** |
| | (0.074) | (0.072) | | (0.081) |
| Manufacturing Employment Decline, 1970–2020 | 1.122* | 1.064 | | 1.069 |
| | (0.063) | (0.065) | | (0.069) |
| Populist Support for Trump | 0.768*** | 0.829*** | | 0.827*** |
| | (0.042) | (0.045) | | (0.052) |
| JBS Membership (log, z) | | 1.504*** | 1.646*** | 1.443*** |
| | | (0.104) | (0.116) | (0.141) |
| House Rep Objected | | | | 1.007 |
| | | | | (0.129) |
| JBS × Objected | | | | 1.110 |
| | | | | (0.108) |
| *Controls* | | | | |
| % Non-Hispanic White | 1.215** | 1.161* | 1.048 | 1.167* |
| | (0.078) | (0.076) | (0.064) | (0.086) |
| Urban County | 2.882*** | 2.639*** | 3.423*** | 2.653*** |
| | (0.384) | (0.351) | (0.432) | (0.396) |
| Distance to DC | 0.739*** | 0.692*** | 0.697*** | 0.694*** |
| | (0.046) | (0.043) | (0.041) | (0.049) |
| County Population | 1.693*** | 1.411*** | 1.420*** | 1.421*** |
| | (0.122) | (0.095) | (0.108) | (0.110) |
| $N$ | 3,112 | 3,112 | 3,112 | 3,069 |
| AIC | 3,128.7 | 3,083.2 | 3,131.7 | 3,042.3 |
| Log-Likelihood | $-1,556.4$ | $-1,532.6$ | $-1,559.9$ | $-1,510.1$ |

*Notes:* Coefficients are incidence rate ratios (IRR) from negative binomial regressions, with HC1 robust standard errors on the IRR scale in parentheses. All continuous covariates are standardized to mean zero and standard deviation one within each estimation sample. IRR $> 1$ indicates a higher defendant count; IRR $< 1$ indicates a lower defendant count. JBS Membership is log(1 + count), standardized. Models including the certification vote are estimated on the 3,069 counties with non-missing House-vote data; the published counterpart has $N = 3{,}071$, with the two-county difference reflecting district–county assignment. Following Pape, Larson, and Ruby (2024), estimation excludes Alaska, which does not report presidential returns by county equivalent. $^{*}p < 0.05$; $^{**}p < 0.01$; $^{***}p < 0.001$.

Table 6: Published and replicated January 6 coefficients

| Variable | Published Model 4 | Replication Model 1 |
|---|---|---|
| White population decline | 1.37 | 1.368 (0.074) |
| Manufacturing employment decline | 1.12 | 1.122 (0.063) |
| Populist Trump support | 0.77 | 0.768 (0.042) |
| % non-Hispanic white | 1.21 | 1.215 (0.078) |
| Urban county | 2.88 | 2.882 (0.384) |
| Distance to Washington, DC | 0.74 | 0.739 (0.046) |
| County population | 1.69 | 1.693 (0.122) |
| $N$ | 3,112 | 3,112 |

*Notes:* Entries are incidence rate ratios. Published values are from Pape et al. (2024, Table 2, Model 4). Replication values are Model 1 from Table 5, with HC1 robust standard errors shown in parentheses.

Table 7: Pape outbidding interaction and JBS membership

| | Published Model 9 | M5a replication | M5b + JBS |
|---|---|---|---|
| White population decline | 1.28$^{***}$ | 1.350$^{***}$ | 1.287$^{**}$ |
| House Republican objected | 0.92 | 0.948 | 1.012 |
| White decline × objected | 1.25$^{*}$ | 1.192$^{\sim}$ | 1.184$^{\sim}$ |
| JBS membership (log, z) | — | — | 1.604$^{***}$ |
| Interaction $p$-value | — | 0.078 | 0.094 |
| $N$ | 3,071 | 3,069 | 3,069 |

*Notes:* Entries are incidence rate ratios. Published values are from Pape et al. (2024, Table 3, Model 9). M5a reproduces the outbidding interaction using the corrected non-missing House-vote sample; M5b adds standardized logged JBS membership. The robust standard error for the JBS coefficient in M5b is 0.112. The outbidding interaction replicates in sign and approximate magnitude and remains marginal after adding JBS membership, while the JBS coefficient enters strongly. $^{\sim}p < 0.10$; $^{*}p < 0.05$; $^{**}p < 0.01$; $^{***}p < 0.001$.

Table 8: Manufacturing decline before and after adding JBS membership

| | Without JBS (Model 1) | | With JBS (Model 2) | |
|---|---|---|---|---|
| Variable | IRR | *p*-value | IRR | *p*-value |
| White Population Decline | 1.368 | $< 0.001$ | 1.332 | $1.3 \times 10^{-7}$ |
| Manufacturing Decline | 1.122 | 0.040 | 1.064 | 0.308 |
| Populist Trump Support | 0.768 | $< 0.001$ | 0.829 | $< 0.001$ |
| **JBS Membership** | — | — | **1.504** | $< \mathbf{0.001}$ |
| AIC | 3128.7 | | 3083.2 | |
| ΔAIC | — | | −45.6 | |

*Notes:* Entries are incidence rate ratios from the January 6 negative binomial specification. Both models include controls for % non-Hispanic white, urban status, distance to Washington, DC, and county population. Standard errors are HC1 robust.

Table 9: Trump populist vote gains and JBS membership (Broz et al. 2021 replication)

| | (1) Replication | (2) + JBS (log, z) | (3) + Any JBS | (4) + JBS p.c. | (5) + 2012 baseline |
|---|---|---|---|---|---|
| JBS membership (log, z) | | $-0.912^{***}$ | | | $-0.759^{***}$ |
| | | (0.162) | | | (0.136) |
| Any JBS (1=yes) | | | $0.494^{\sim}$ | | |
| | | | (0.295) | | |
| JBS per 10,000 (log, z) | | | | $-0.786^{***}$ | |
| | | | | (0.146) | |
| Romney 2012 share | | | | | $-0.148^{***}$ |
| | | | | | (0.008) |
| Mfg. decline 1970–2015 | $0.045^{**}$ | $0.040^{**}$ | $0.045^{**}$ | $0.040^{**}$ | 0.019 |
| | (0.015) | (0.014) | (0.015) | (0.014) | (0.012) |
| Log population (2015) | $-0.284^{\sim}$ | $0.355^{\sim}$ | $-0.303^{\sim}$ | $-0.297^{*}$ | 0.004 |
| | (0.159) | (0.187) | (0.161) | (0.148) | (0.168) |
| Additional demographic controls | Yes | Yes | Yes | Yes | Yes |
| *N* | 2,608 | 2,608 | 2,608 | 2,608 | 2,608 |

*Notes:* Population-weighted OLS with HC1 robust standard errors in parentheses, following Broz et al. (2021, Table A1). Dependent variable: Trump 2016 two-party vote share minus Romney 2012. All columns include the additional Broz et al. demographic controls omitted from the table: % age 65+, % white, % Hispanic, log wage, % female, % college, and business/services share. Column (5) additionally includes Romney's 2012 county two-party share. The sign change on log population between (1) and (2) reflects collinearity between membership counts and population; other coefficients are stable. Dropping California ($\beta = -0.78^{***}$) or Los Angeles County ($\beta = -0.94^{***}$) leaves the JBS estimate essentially unchanged. $^{\sim}p < 0.10$; $^{*}p < 0.05$; $^{**}p < 0.01$; $^{***}p < 0.001$.

Table 10: Right-wing terrorism logits and JBS membership (Nemeth & Hansen 2021 replication)

| Specification | Coefficient | SE | $p$ | $N$ |
|---|---|---|---|---|
| Main model: JBS membership (log, z) | -0.020 | 0.141 | 0.910 | 133,539 |
| Any JBS indicator | -0.174 | 0.411 | 0.672 | 133,539 |
| JBS × electoral competition | -0.229 | 0.287 | 0.420 | 133,539 |
| JBS × racial composition | -0.006 | 0.005 | 0.260 | 133,539 |
| 1988–2016 window: JBS membership | -0.211 | 0.180 | 0.240 | 85,457 |
| 1971–1987 window: JBS membership | 0.238 | 0.251 | 0.343 | 48,082 |

*Notes:* Entries are logit coefficients. All specifications include the Nemeth & Hansen control vector: lagged political competition, racial composition, county GDP per capita, labor-force participation, citizenship, population, urbanization, land area, spatial lag of right-wing violence, election-year indicators, presidential-party controls, region indicators, and cubic time trends. Standard errors are clustered by county. The 1988–2016 and 1971–1987 windows use the corrected calendar-year variable from the Stata date field.

Table 11: Sensitivity of January 6 models to the population control

| Specification | JBS term | JBS IRR | $p$ | AIC |
|---|---|---|---|---|
| Published population form | log JBS, standardized | 1.504 | $\approx 4 \times 10^{-9}$ | 3,083.2 |
| Log-population control | log JBS, standardized | 0.930 | 0.280 | 2,900.4 |
| Log-population control | any JBS indicator | 0.800 | 0.189 | 2,900.6 |

*Notes:* Entries are incidence rate ratios from negative binomial specifications with HC1 robust standard errors. The log-population versions replace Pape et al.'s linear county-population control, measured in 100,000s, with standardized log population.

Table 12: Average JBS membership by congressional-district certification vote

| | Objected | Certified | Ratio | $p$ |
|---|---|---|---|---|
| All JBS members | 117.2 | 93.3 | 1.26 | 0.002 |
| *By gender:* | | | | |
| Male | 68.7 | 54.4 | 1.26 | 0.001 |
| Female | 41.8 | 34.1 | 1.22 | 0.007 |
| *By chapter type:* | | | | |
| Local chapter | 73.7 | 55.4 | 1.33 | 0.005 |
| HOME (national) | 41.4 | 36.2 | 1.14 | 0.002 |

*Notes:* Unit is the congressional district. Entries are average JBS members per district, comparing the 139 Republican objector districts with the 294 occupied districts whose representatives certified the result. The universe is the 433 seats occupied on January 6, 2021; LA–05 was vacant and NY–22 was unresolved. $p$-values are from one-sided Mann–Whitney $U$ tests, with zero-JBS districts included.

## 7.7 Supplementary Diagnostic and Conceptual Figures

The appendix figure clarifies the capacity-versus-conversion interpretation behind the sign reversal between the January 6 and Broz electoral-conversion results.

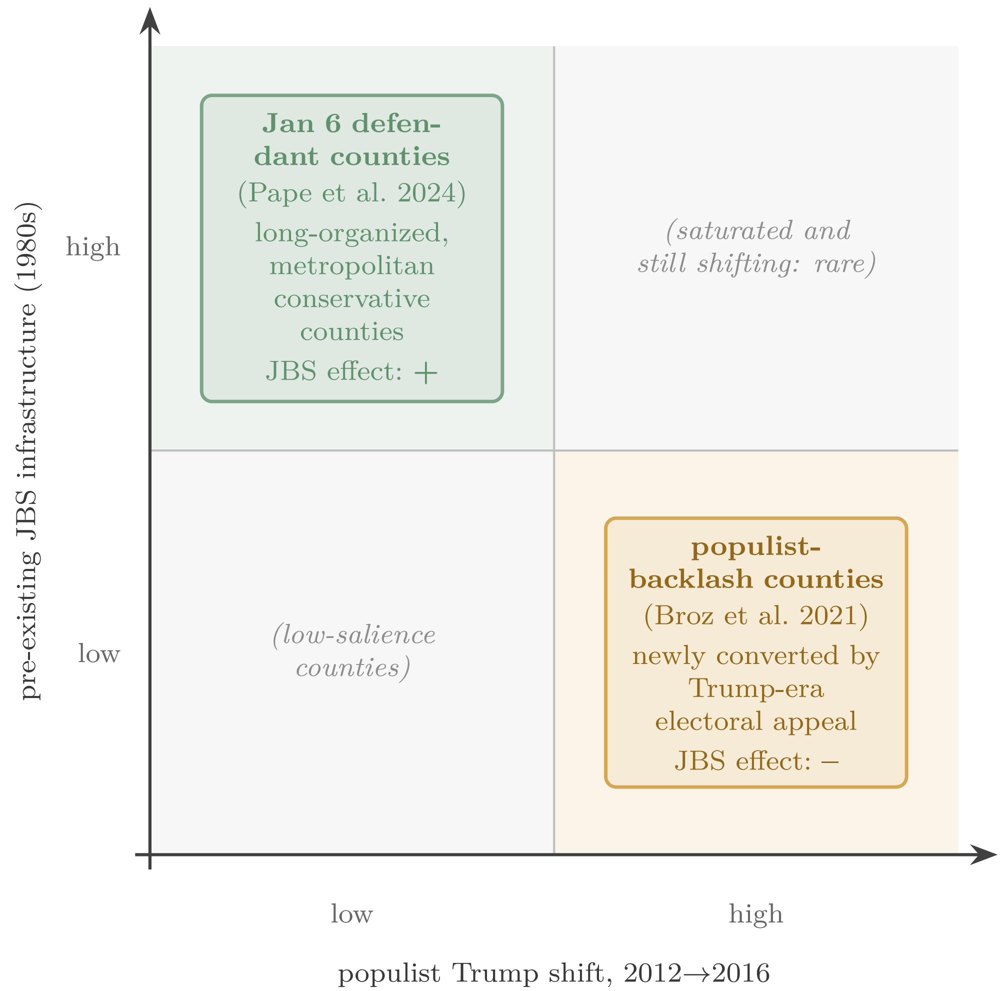


Figure 10: A capacity-versus-conversion interpretation of the sign reversal. JBS-rich counties combine high mobilizing capacity with low 2012–2016 electoral conversion; controlling for the 2012 Republican baseline attenuates the JBS coefficient only modestly ($-0.91 \rightarrow -0.76$), indicating that the pattern reflects the metropolitan geography of older conservative organization rather than a ceiling effect in already-Republican counties.

# References


Acharya, Avidit, Matthew Blackwell, and Maya Sen (July 2016). "The Political Legacy of American Slavery". *The Journal of Politics* 78.3, pp. 621–641. DOI: 10.1086/686631.

Broz, J. Lawrence, Jeffry Frieden, and Stephen Weymouth (Apr. 2021). "Populism in Place: The Economic Geography of the Globalization Backlash". *International Organization* 75.2, pp. 464–494. DOI: 10.1017/S0020818320000314.

Critchlow, Donald T. (2005). *Phyllis Schlafly and Grassroots Conservatism: A Woman's Crusade*. Princeton, NJ: Princeton University Press.

Dallek, Matthew (2023). *Birchers: How the John Birch Society radicalized the American right*. Basic Books.

Dell, Melissa (Mar. 2025). "Deep Learning for Economists". *Journal of Economic Literature* 63.1, pp. 5–58. DOI: 10.1257/jel.20241733.

Dell, Melissa et al. (2023). "American Stories: A Large-Scale Structured Text Dataset of Historical U.S. Newspapers". *Advances in Neural Information Processing Systems*.

DePue, Mark (2011). *Interview with Phyllis Schlafly*. Abraham Lincoln Presidential Library Oral History Program, ISE-A-L-2011-001.01.

Fryer Jr., Roland G. and Steven D. Levitt (Nov. 2012). "Hatred and Profits: Under the Hood of the Ku Klux Klan". *The Quarterly Journal of Economics* 127.4, pp. 1883–1925. DOI: 10.1093/qje/qjs028.

Jocher, Glenn, Ayush Chaurasia, and Jing Qiu (2023). *Ultralytics YOLOv8*. Version 8.0.0. URL: https://github.com/ultralytics/ultralytics.

Kraft, Charles Jeffrey (1992). "A Preliminary Socio-Economic and State Demographic Profile of the John Birch Society". Undergraduate thesis, Tufts University.

Lazar, Ernie (2020). *Schlafly, Phyllis and John Birch Society*. Archival documents including correspondence. URL: archive.org (visited on 03/16/2025).

Li, Minghao et al. (2023). "TrOCR: Transformer-Based Optical Character Recognition with Pre-trained Models". *Proceedings of the AAAI Conference on Artificial Intelligence* 37.11, pp. 13094–13102.

Lin, Tsung-Yi et al. (2014). "Microsoft COCO: Common Objects in Context". In: *European Conference on Computer Vision*. Springer, pp. 740–755.

Madestam, Andreas et al. (Nov. 2013). "Do Political Protests Matter? Evidence from the Tea Party Movement". *The Quarterly Journal of Economics* 128.4, pp. 1633–1685. DOI: 10.1093/qje/qjt021.

McGirr, Lisa (2001). *Suburban Warriors: The Origins of the New American Right*. Princeton, NJ: Princeton University Press.

McVeigh, Rory, David Cunningham, and Justin Farrell (Dec. 2014). "Political Polarization as a Social Movement Outcome: 1960s Klan Activism and Its Enduring Impact on Political Realignment in Southern Counties, 1960 to 2000". *American Sociological Review* 79.6, pp. 1144–1171. DOI: 10.1177/0003122414555885.

Nemeth, Stephen C. and Holley E. Hansen (June 2022). "Political Competition and Right-Wing Terrorism: A County-Level Analysis of the United States". *Political Research Quarterly* 75.2, pp. 338–352. DOI: 10.1177/10659129211006791.

Pape, Robert A., Kyle D. Larson, and Keven G. Ruby (July 2024). "The Political Geography of the January 6 Insurrectionists". *PS: Political Science & Politics* 57.3, pp. 329–339. DOI: 10.1017/S1049096524000040.

Satyanath, Shanker, Nico Voigtländer, and Hans-Joachim Voth (Apr. 2017). "Bowling for Fascism: Social Capital and the Rise of the Nazi Party". *Journal of Political Economy* 125.2, pp. 478–526. DOI: 10.1086/690949.

Spruill, Marjorie (2017). *Divided we stand: The battle over women's rights and family values that polarized American politics*. Bloomsbury Publishing.

Tkachenko, Maxim et al. (2020). *Label Studio*. https://github.com/heartexlabs/label-studio.

Voigtländer, Nico and Hans-Joachim Voth (Aug. 2012). "Persecution Perpetuated: The Medieval Origins of Anti-Semitic Violence in Nazi Germany". *The Quarterly Journal of Economics* 127.3, pp. 1339–1392. DOI: 10.1093/qje/qjs019.

Wu, Yuxin et al. (2019). *Detectron2*. https://github.com/facebookresearch/detectron2.